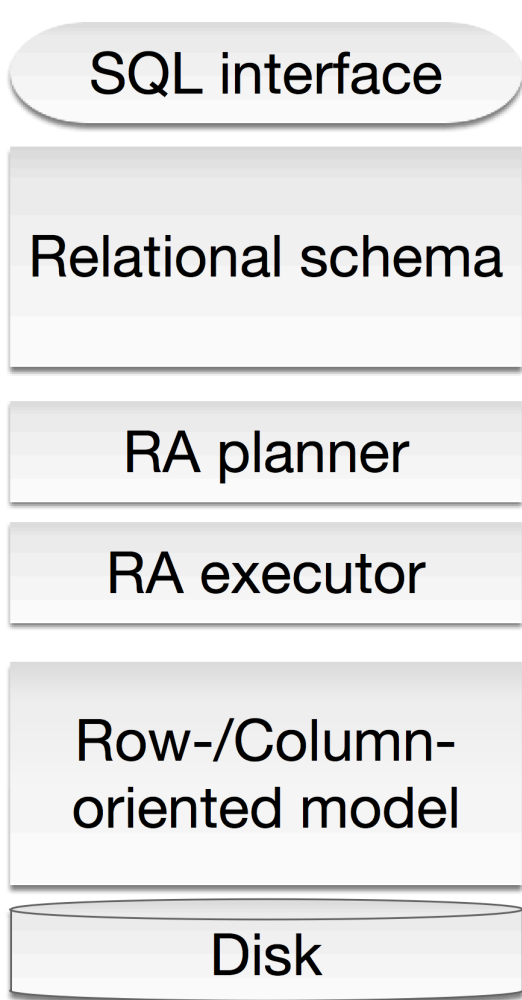


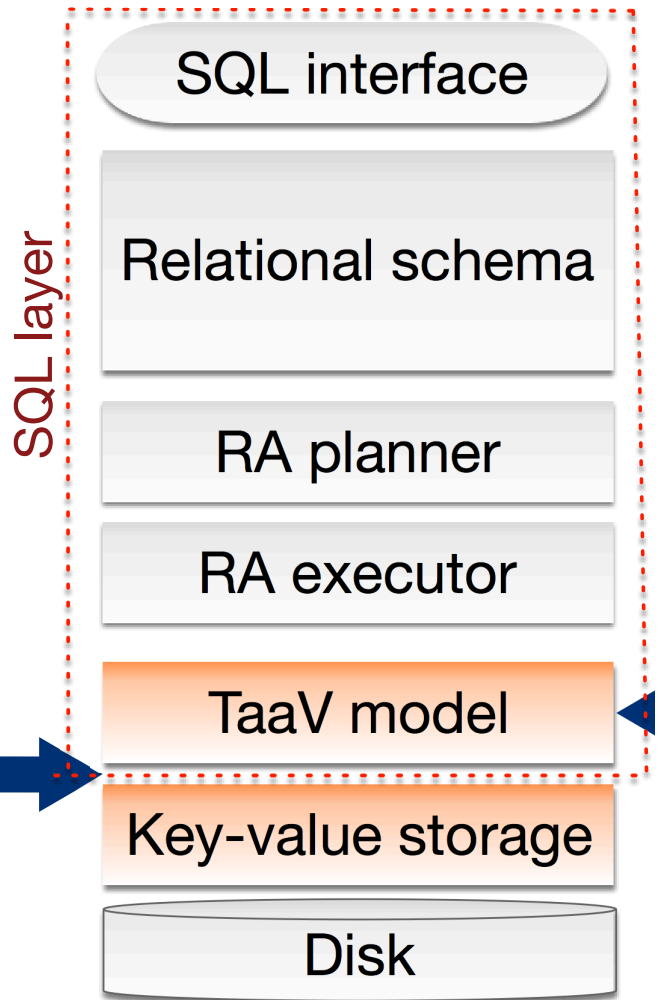
- ▶ Parallel data systems are typically built on key-value storage
- ▶ Following an “SQL-on-KV” paradigm



The SQL@KV-store Paradigm



(a) RDBMS



(b) SQL layers@KV

Tuple-as-a-Value: (key, tuple)

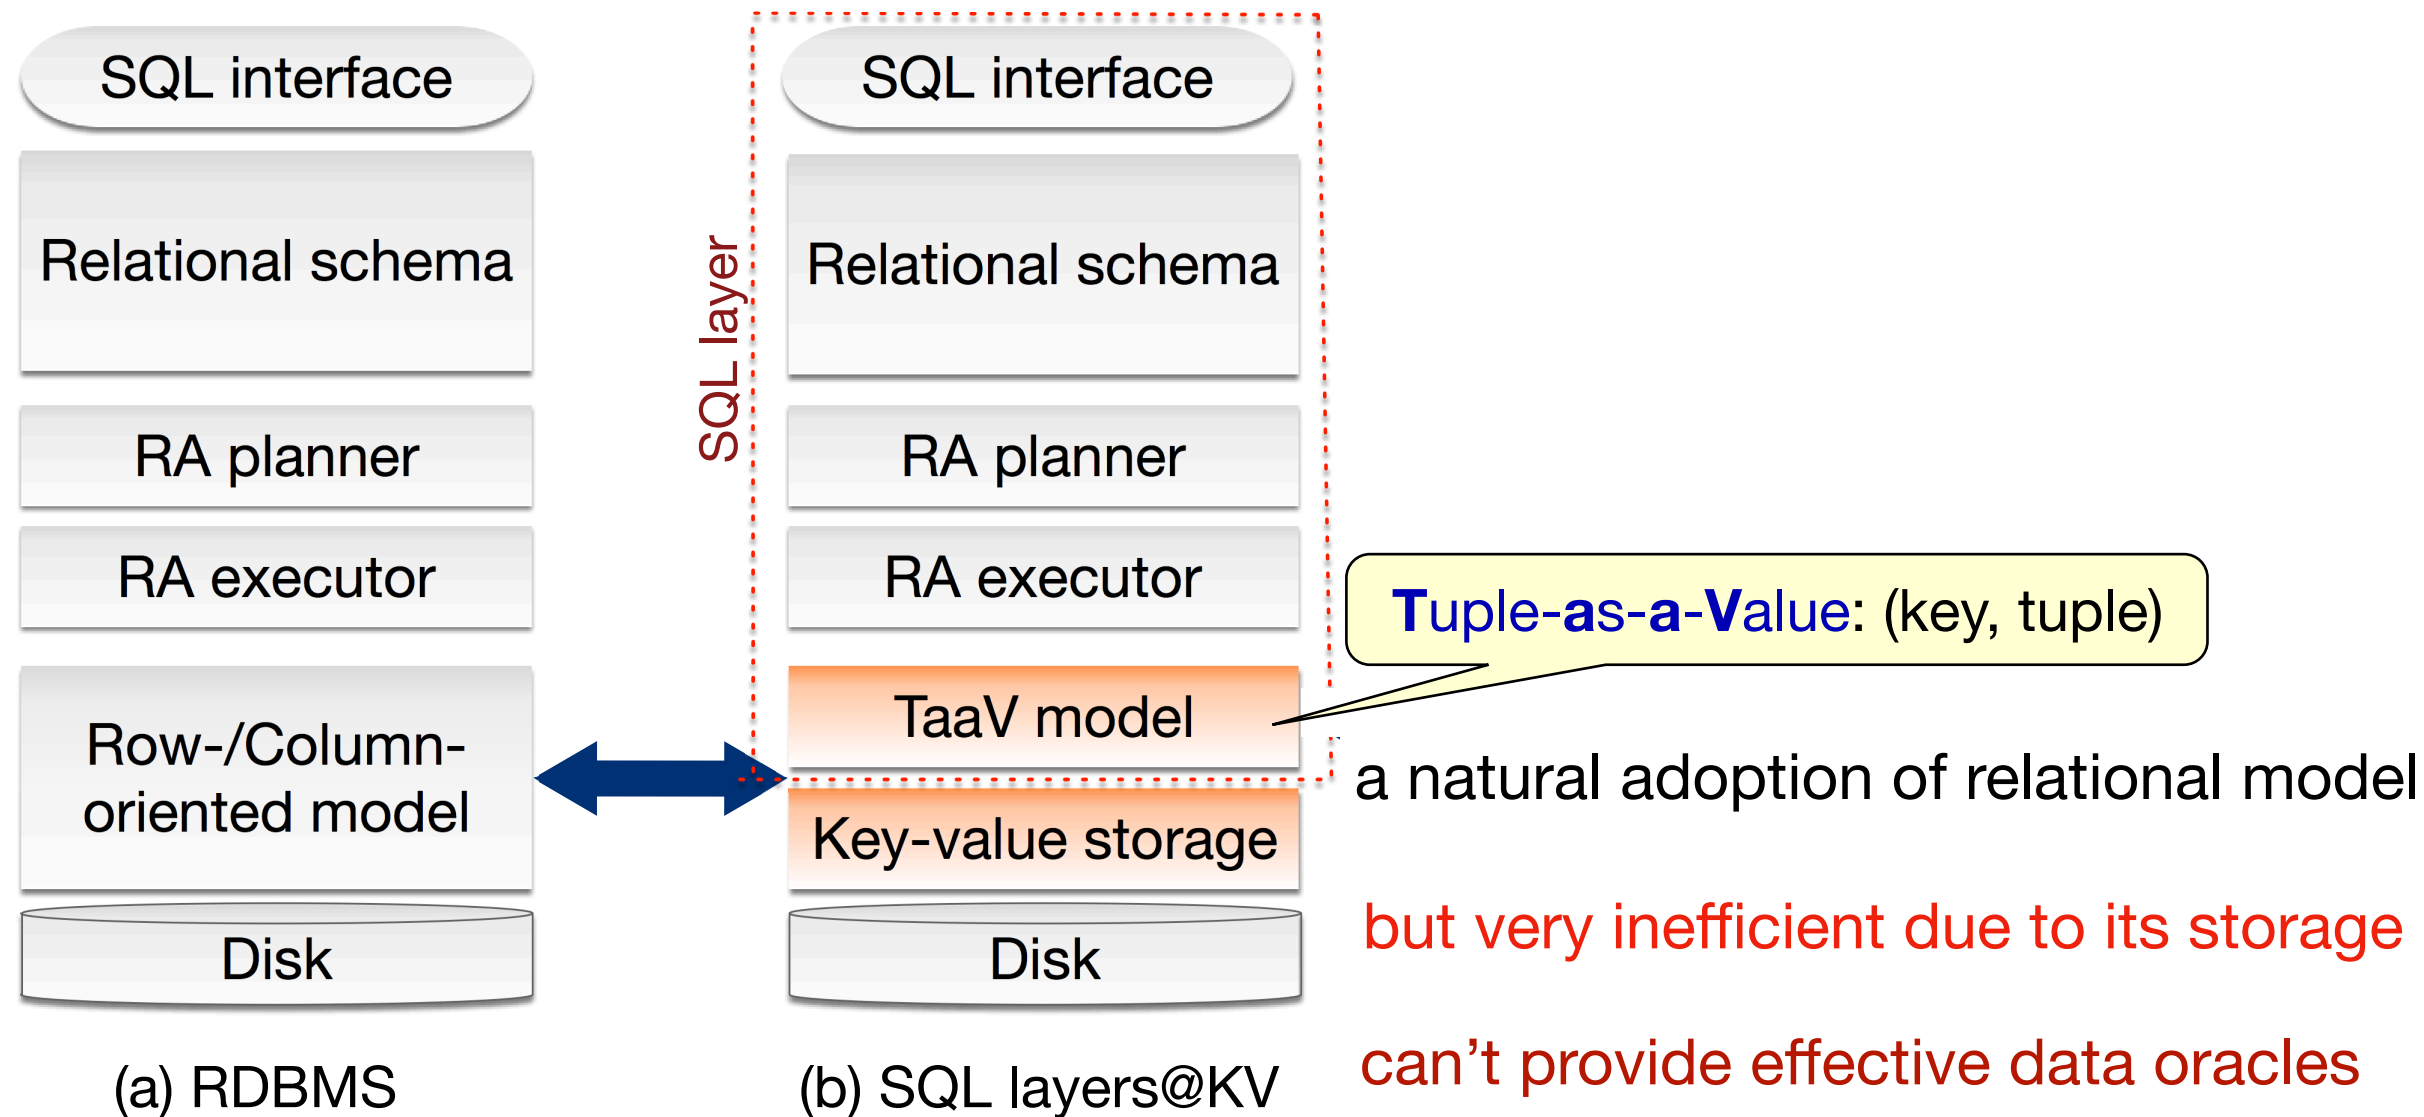
anatural adoption of relational model

but very inefficient due to its storage

can't provide effective data oracles

The SQL@KV-store Paradigm

- ▶ Parallel data systems are typically built on key-value storage
- ▶ Following an “SQL-on-KV” paradigm



The SQL@KV-store Paradigm

- ▶ A new data model for “SQL-on-KV”
 - from TaaV to SaaV (Set-as-a-Value)

