





```
SELECT fid  
FROM    friend AS f, update AS u  
WHERE   f.uid= uid1 AND f.fid=u.uid AND u.country=UK
```

# *A Conventional DBMS plan*

*find all my friends who have online updates read from UK*

- `friend(uid, fid)`
- `update(id, uid, country, date, post, ...)`

$S_1$ : for each friend, find a set  $F_1$  of all my friends

\$2: for each  $f_i$  in  $F_1$ , check whether  $f_i$  has a record in  $update$  from  $U$





*a nested loop*



# A Conventional Query Plan



*a linear scan*

# Facebook:

- ▶ **billions** of friend tuples
- ▶ **trillions** of update tuples



Fast! 12GB/s

► S1 takes 20 mins ( $|F1|=100$ )

► S2 takes 2000 mins

*Polynomial time queries become intractable on big data*



# A Conventional Query Plan

*find all my friends who have online update records from UK*

- **friend**(uid, fid)
- **update**(id, uid, country, date, post, ...)

```
SELECT fid
FROM   friend AS f, update AS u
WHERE  f.uid= uid1 AND f.fid=u.uid AND u.country=UK
```

*a linear scan*

*A Conventional DBMS plan*

S1: for each **friend**, find a set  $F_1$  of all my friends

*a nested loop*

S2: for each fid  $f_0$  in  $F_1$ , check whether  $f_0$  has a record in **update** from UK

Facebook:

► **billions** of **friend** tuples




► S1 takes **20 mins** ( $|F_1|=100$ )

*Polynomial time queries become intractable on big data*

# Common Wisdom

## ► Parallel query processing



Assuming linear scalability, using 50,000 processors

- 2,020 mins is reduced to 2.4s