



Bounded Approximation Scheme

Input: A resource ratio $\alpha \in (0, 1]$ and an access schema \mathcal{A} .

Scheme^[6]: Given generic SQL Q and D , computes $(Q(D_Q), \boldsymbol{\eta})$:

- ▶ access a fraction D_Q of D with $|D_Q| \leq \alpha |D|$;
- ▶ $\text{accuracy}(Q, D, Q(D_Q)) \geq \eta$.

Challenges:

- ▶ **Deterministic** bound for **generic** queries (even non-aggregate)
- ▶ Compute **both answers** and **accuracy bound** by **accessing $\leq \alpha|D|$ tuples**

decided by available resources

approximate answers

A yellow speech bubble with a black outline, containing the text 'approximate answers'. The bubble has a rectangular body with rounded corners and a triangular tail pointing downwards and to the right.



accuracy bound



deterministic!

Flexible trade-offs: available resource vs. accuracy bound

[6] Y. Cao, W. Fan: Data Driven Approximation with Bounded Resources, VLDB 2017


not *one-size-fits-all* reduction





accesses α/δ tuples





Bounded Approximation Scheme

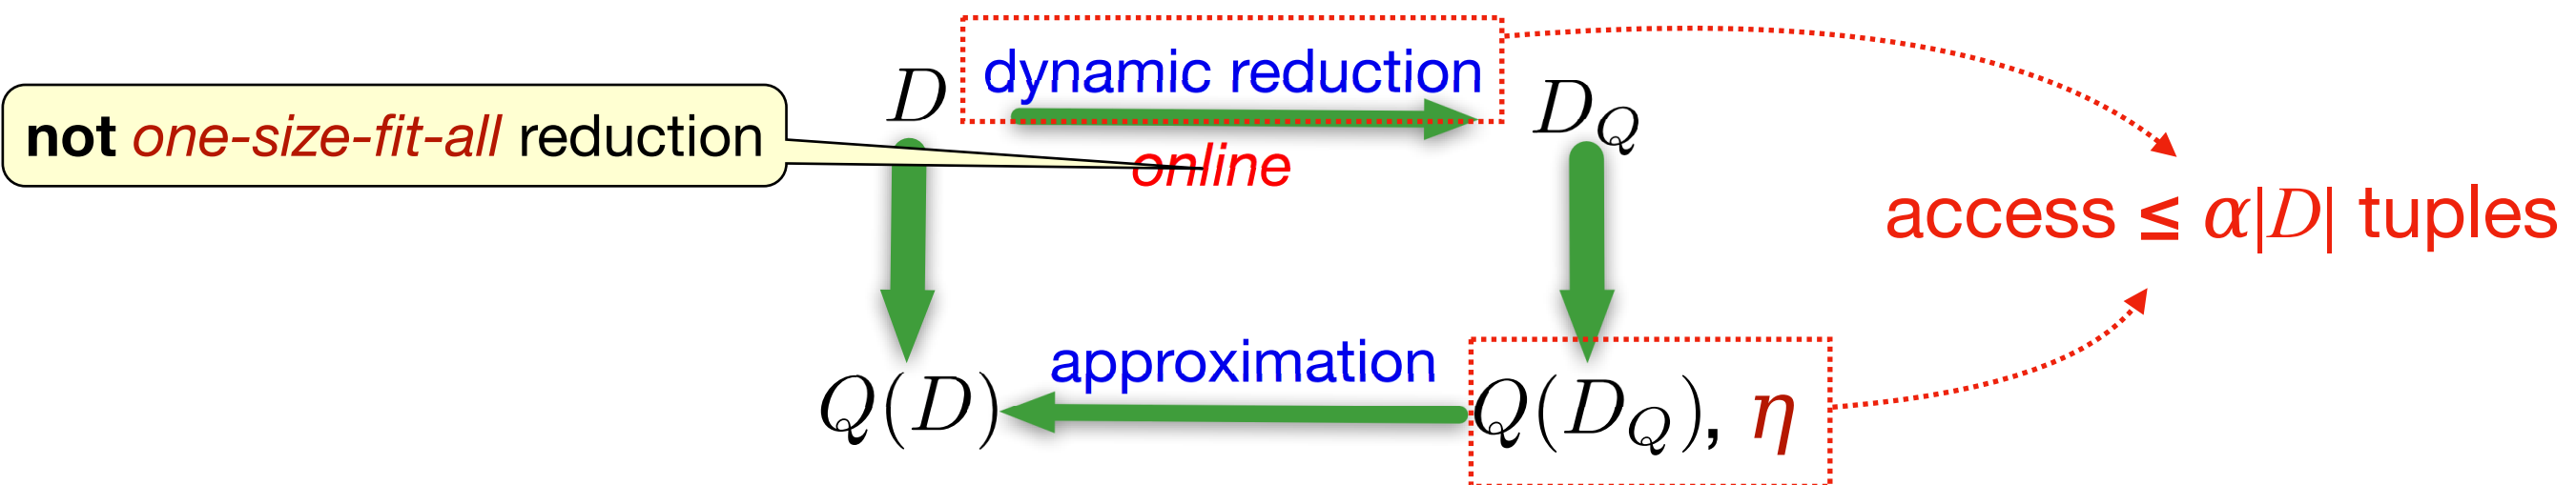
decided by available resources

approximate answers

Input: A resource ratio $\alpha \in (0, 1]$ and an access schema \mathcal{A} .

Scheme^[6]: Given generic SQL Q and D , computes $(Q(D_Q), \eta)$:

- ▶ access a fraction D_Q of D with $|D_Q| \leq \alpha|D|$; accuracy bound
- ▶ $\text{accuracy}(Q, D, Q(D_Q)) \geq \eta$. deterministic!



Flexible trade-offs: available resources vs. accuracy bound

Challenges:

- ▶ **Deterministic** bound for **generic** queries (even non-aggregate)
- ▶ Compute **both** answers and accuracy bound by **accessing** $\leq \alpha|D|$ tuples

How does it work?
